

# Topic 11 – Multiple Linear Regression

ENVX1002 Introduction to Statistical Methods

**Liana Pozza**

*The University of Sydney*

Dec 2024



THE UNIVERSITY OF  
**SYDNEY**

# Module overview

- **Week 9. Describing relationships**

- ➡ Correlation → calculation, interpretation, things to watch out for
- ➡ Regression → Why do we care? model structure, model fitting

- **Week 10. Linear functions**

- ➡ Is the model worth fitting? → Assumptions, hypothesis testing
- ➡ How good is the model? → Measures of model fit

- **Week 11. Linear functions - multiple predictors**

- ➡ Parsimonious models
- ➡ Introduction to Multiple Linear Regression (MLR) modelling
- ➡ Assumptions and interpretation

- **Week 12. Nonlinear functions**

- ➡ Common nonlinear functions
- ➡ Transformations
- ➡ Performing nonlinear regression

# Module overview

- **Week 11. Linear functions - multiple predictors**
  - ➡ Parsimonious models
  - ➡ Introduction to Multiple Linear Regression (MLR) modelling
  - ➡ Assumptions and interpretation

# Recap

# Simple linear regression


$$Y_i = \beta_0 + \beta_1 x_i + \epsilon_i$$

Ideal for predicting a continuous response variable from a single predictor variable: *“How does  $y$  change as  $x$  changes?”*

## What if we have more than one predictor?

What is the model and how do we interpret the results?

# Multiple linear regression



*Francis Galton and Karl Pearson*

# History

- First raised by **Francis Galton** in 1886, after studying genetic variations in sweet peas over several generations.
- **Karl Pearson** developed the mathematical formalism for the multiple linear regression model in the early 1900s.

*“The somewhat complicated mathematics of multiple correlation, with its repeated appeals to the geometrical notions of hyperspace, remained a closed chamber to him.”*

– Pearson (1930), on Galton’s work with MLR

# Air Quality in New York (1973)



# Air quality

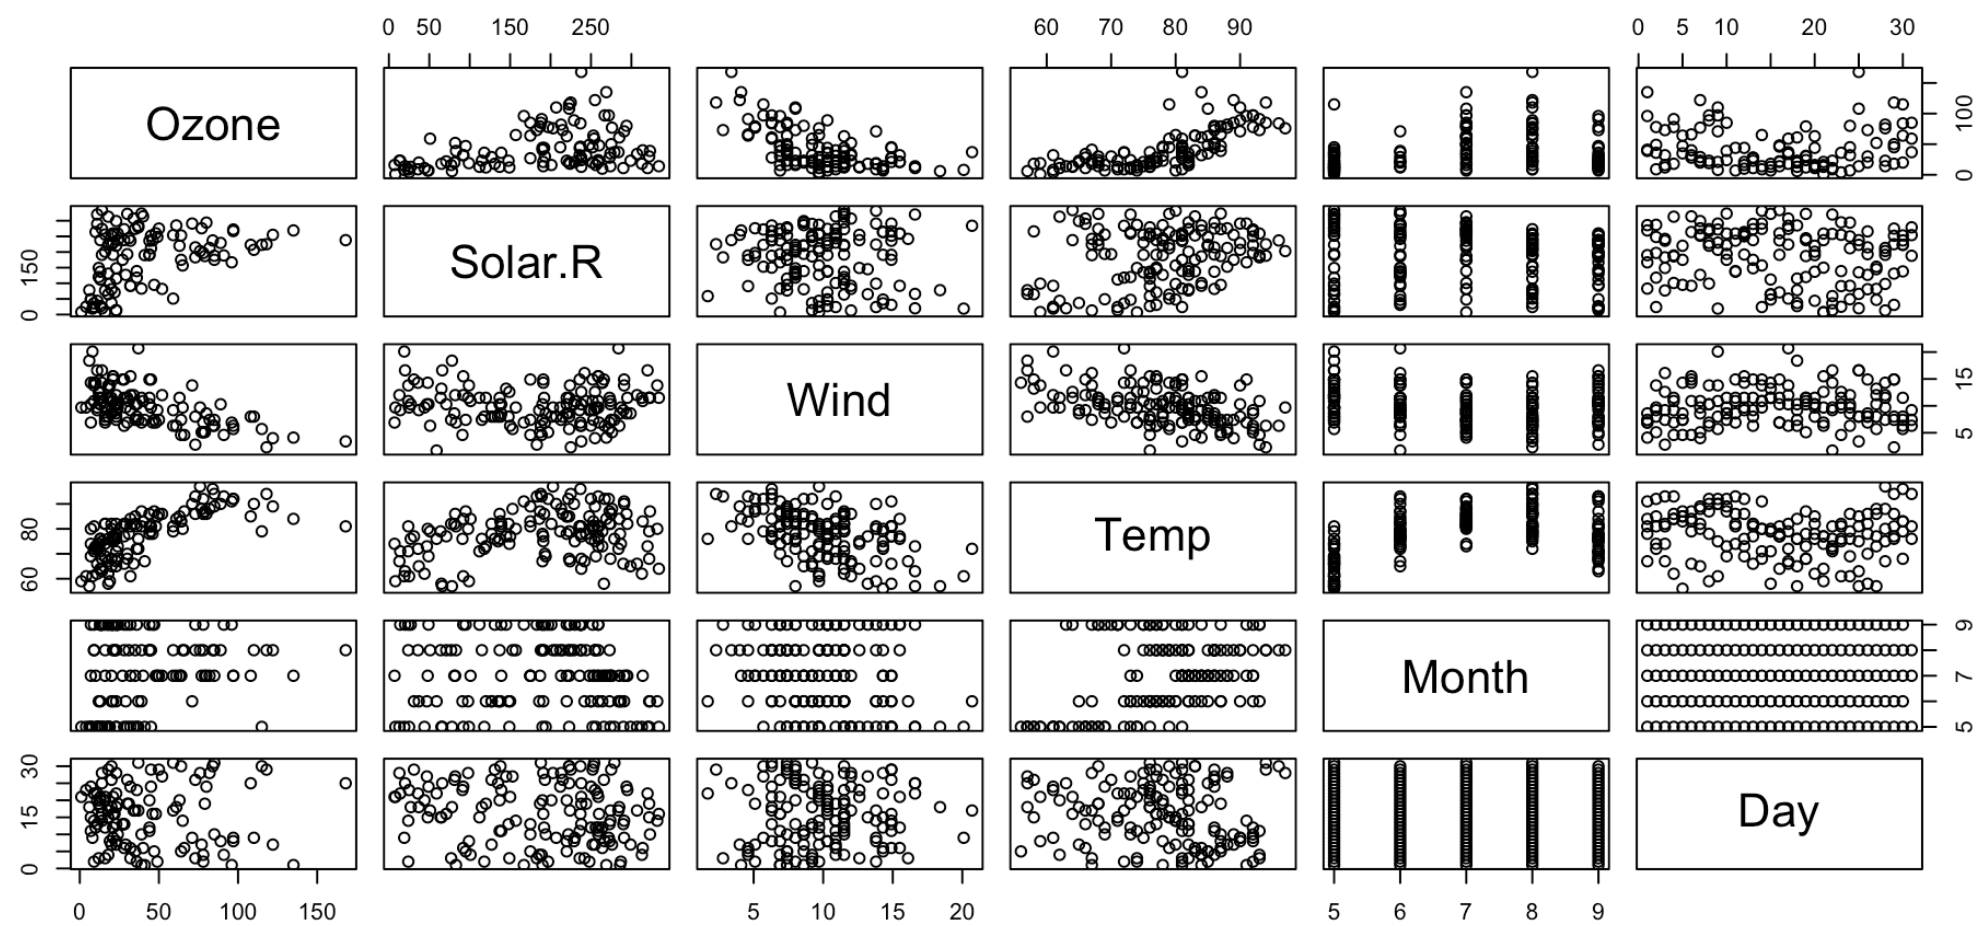
```
Rows: 153
Columns: 6
$ Ozone    <int> 41, 36, 12, 18, NA, 28, 23, 19, 8, NA, 7, 16, 11, 14, 18, 14, ...
$ Solar.R  <int> 190, 118, 149, 313, NA, NA, 299, 99, 19, 194, NA, 256, 290, 27...
$ Wind     <dbl> 7.4, 8.0, 12.6, 11.5, 14.3, 14.9, 8.6, 13.8, 20.1, 8.6, 6.9, 9...
$ Temp     <int> 67, 72, 74, 62, 56, 66, 65, 59, 61, 69, 74, 69, 66, 68, 58, 64...
$ Month    <int> 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, ...
$ Day      <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, ...
```

```
Rows: 153
Columns: 6
$ Ozone    <int> 41, 36, 12, 18, NA, 28, 23, 19, 8, NA, 7, 16, 11, 14, 18, 14, ...
$ Solar.R  <int> 190, 118, 149, 313, NA, NA, 299, 99, 19, 194, NA, 256, 290, 27...
$ Wind     <dbl> 7.4, 8.0, 12.6, 11.5, 14.3, 14.9, 8.6, 13.8, 20.1, 8.6, 6.9, 9...
$ Temp     <int> 67, 72, 74, 62, 56, 66, 65, 59, 61, 69, 74, 69, 66, 68, 58, 64...
$ Month    <int> 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, ...
$ Day      <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, ...
```


Ozone: harmful air pollutant when present at ground level; main component of smog:

- `Ozone`: ozone concentration (ppb)
- `Solar.R`: solar radiation (lang)
- `Wind`: wind speed (mph)
- `Temp`: ambient temperature (degrees F)
- `Month`: month (1-12)
- `Day`: day of the month (1-31)

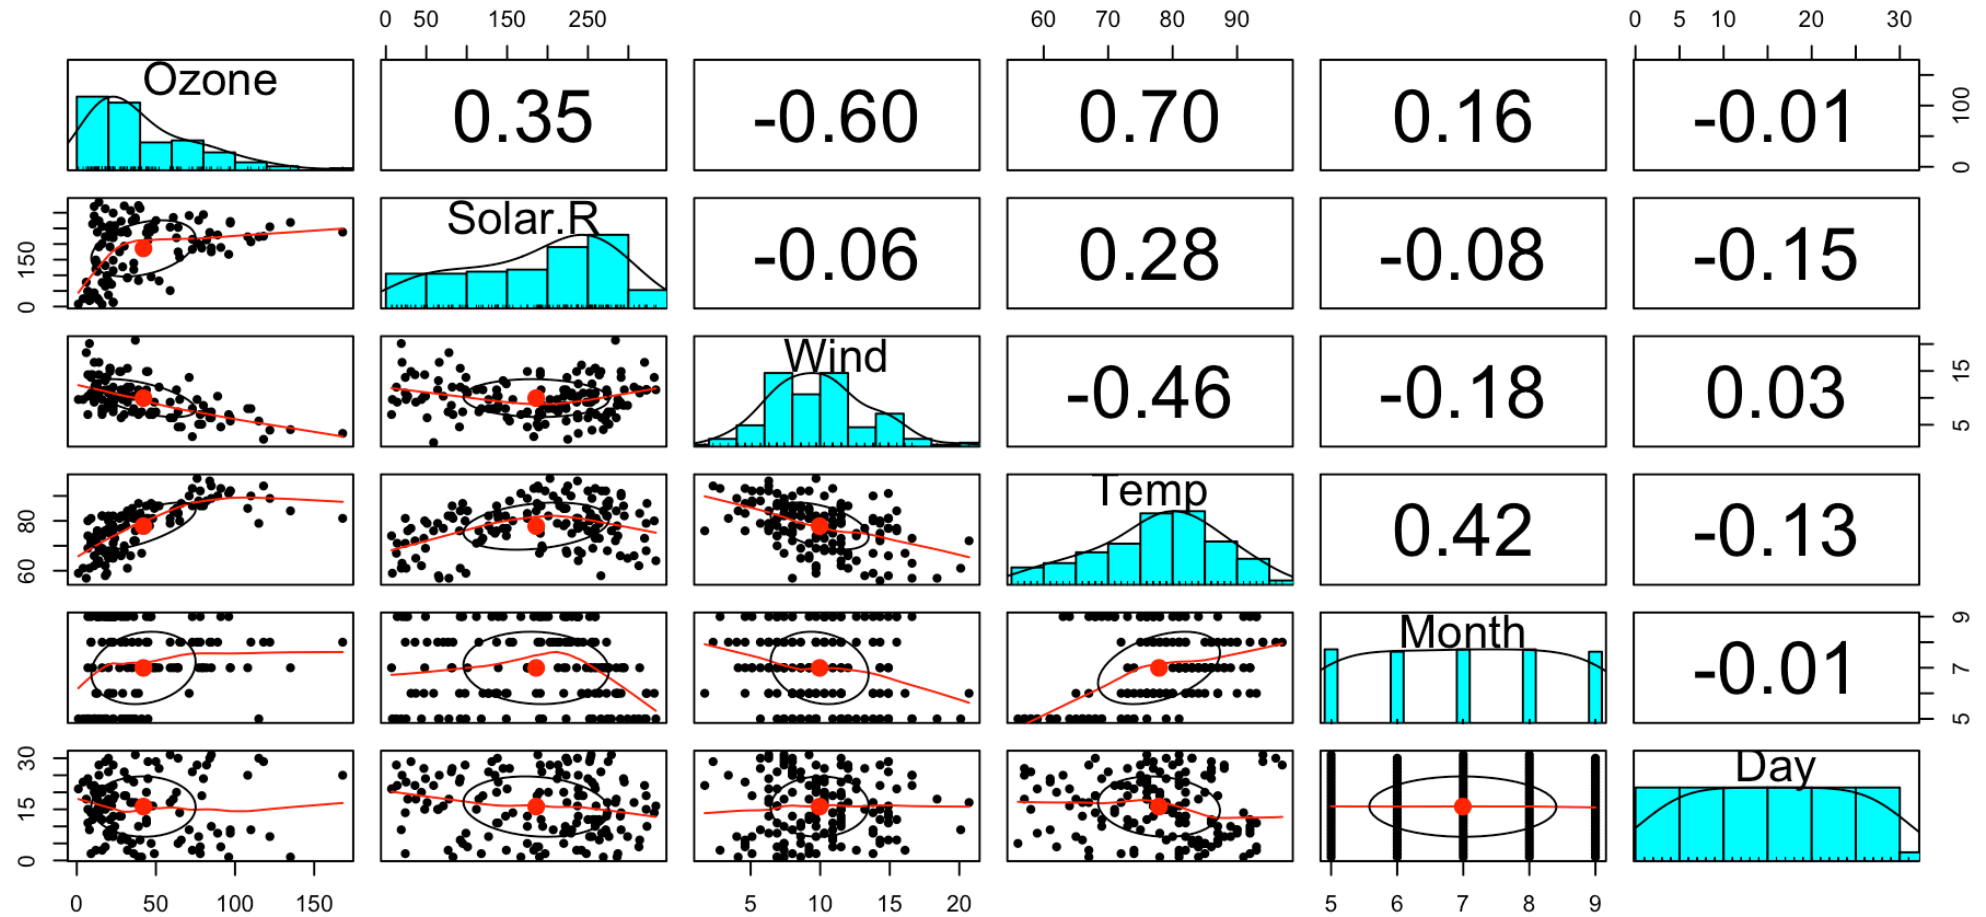
# Correlations



# corrplot



# psych



# The simplest model


Pick the predictor that has the highest correlation coefficient with the response variable.

	Ozone	Solar.R	Wind	Temp	Month
Ozone	1.000000000	0.34834169	-0.61249658	0.6985414	0.142885168
Solar.R	0.348341693	1.00000000	-0.12718345	0.2940876	-0.074066683
Wind	-0.612496576	-0.12718345	1.00000000	-0.4971897	-0.194495804
Temp	0.698541410	0.29408764	-0.49718972	1.0000000	0.403971709
Month	0.142885168	-0.07406668	-0.19449580	0.4039717	1.000000000
Day	-0.005189769	-0.05775380	0.04987102	-0.0965458	-0.009001079
Day					
Ozone	-0.005189769				
Solar.R	-0.057753801				
Wind	0.049871017				
Temp	-0.096545800				
Month	-0.009001079				
Day	1.000000000				

What can we understand about the relationship between Ozone and Temp ( $r = 0.7$ )?


# Relationship

What can we understand about the relationship between Ozone and Temp ( $r = 0.7$ )?



# Relationship





What can we understand about the relationship between Ozone and Temp ( $r = 0.7$ )?




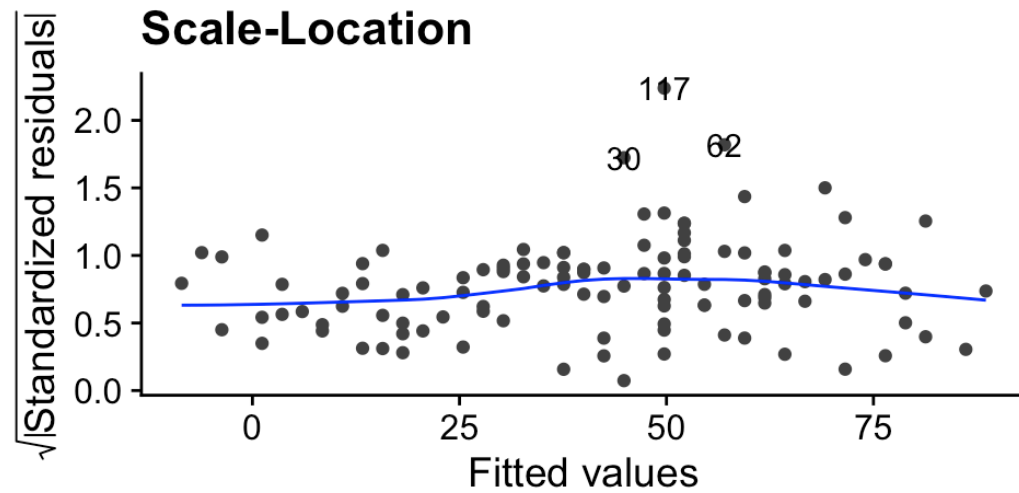

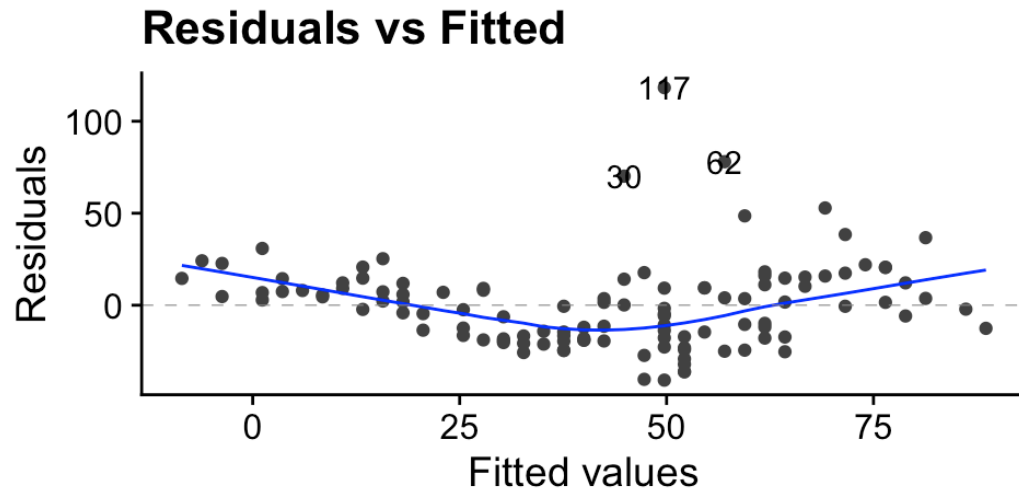


# Fitting the model

# Assumptions



# ggfortify



# Interpretation

```
Call:
lm(formula = Ozone ~ Temp, data = airquality)

Residuals:
    Min       1Q   Median       3Q      Max
-40.729 -17.409  -0.587   11.306  118.271

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -146.9955     18.2872  -8.038 9.37e-13 ***
Temp          2.4287      0.2331   10.418 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 23.71 on 114 degrees of freedom
(37 observations deleted due to missingness)
Multiple R-squared:  0.4877,    Adjusted R-squared:  0.4832
F-statistic: 108.5 on 1 and 114 DF,  p-value: < 2.2e-16
```

- Temp is a statistically significant predictor of Ozone ( $p < .001$ ).
- The (simple linear) model explains 49% of variance ( $r^2 = 0.49$ ).

Can we improve the model in other ways?

# Multiple linear regression

# Important concepts

- The “best” model is the one that best describes the relationship between the response and the predictors.
  - ➡ **NOT** the model that includes all possible predictors ([data dredging](#)).

## Principle of parsimony

A good model:

- Has only *useful* predictors.
- Has no *redundant* predictors (principle of orthogonality).
- Is *interpretable* (principle of transparency) or *predicts well* (principle of accuracy).

# The MLR model

An extension of simple linear regression to include **more than one** predictor variable: “How does ***y*** change as ***x*<sub>1</sub>, *x*<sub>2</sub>, ..., *x*<sub>*k*</sub>** change?”

$$Y_i = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k + \epsilon_i$$

Therefore, estimating the model involves *estimating the values of  $\beta_0, \beta_1, \beta_2, \dots, \beta_k$* .

- $\beta_0$  is the intercept
- $\beta_1$  to  $\beta_k$  are the partial regression coefficients
- $\epsilon$  is the error term

# Explore

```
Rows: 153
Columns: 6
$ Ozone    <int> 41, 36, 12, 18, NA, 28, 23, 19, 8, NA, 7, 16, 11, 14, 18, 14, ...
$ Solar.R  <int> 190, 118, 149, 313, NA, NA, 299, 99, 19, 194, NA, 256, 290, 27...
$ Wind     <dbl> 7.4, 8.0, 12.6, 11.5, 14.3, 14.9, 8.6, 13.8, 20.1, 8.6, 6.9, 9...
$ Temp     <int> 67, 72, 74, 62, 56, 66, 65, 59, 61, 69, 74, 69, 66, 68, 58, 64...
$ Month    <int> 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, ...
$ Day      <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, ...
```

## The “best” model

The variables `Month` and `Day` are not useful predictors, so we will exclude them from the model.



# Visualisation: not easy

Are the plots useful?

## 3D plot

WebGL is not supported by your browser -  
visit <https://get.webgl.org> for more info

# Visualisation: not easy

Are the plots useful?

## 4D plot

WebGL is not supported by your browser -  
visit <https://get.webgl.org> for more info

# Partial regression coefficients

Given the multiple linear model:

$$Y_i = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k + \epsilon_i$$

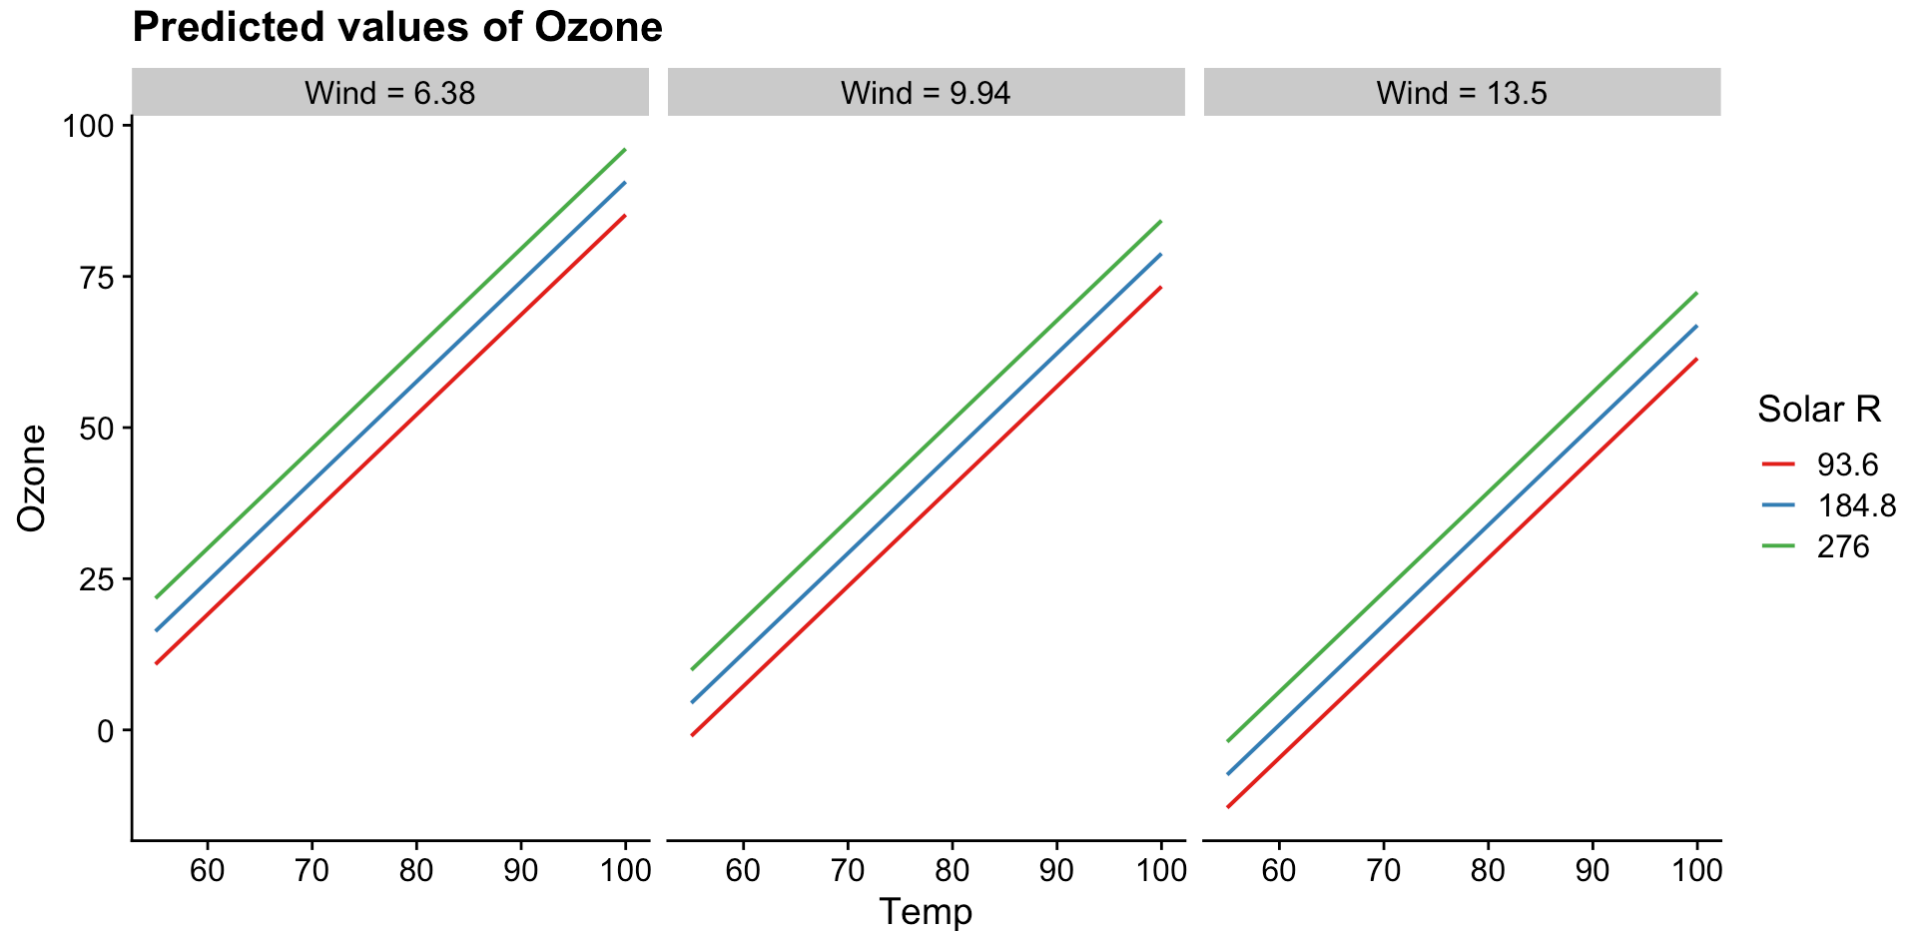
The partial regression coefficient for a predictor  $x_i$  is the amount by which the response variable  $Y$  changes when  $x_k$  is increased by one unit, **while all other predictors are held constant**.

$$\beta_k = \frac{\Delta Y}{\Delta x_k}$$

$$\text{Ozone} = \alpha + \beta_1(\text{Solar.R}) + \beta_2(\text{Wind}) + \beta_3(\text{Temp}) + \epsilon$$

With *Wind* and *Solar.R* held constant, how does *Temp* affect *Ozone*?

# Partial regression coefficients: visualisation



With *Wind* and *Solar.R* held constant, how does *Temp* affect *Ozone*?

# Interpreting the partial regression coefficients

```
Call:
lm(formula = Ozone ~ Solar.R + Wind + Temp, data = airquality)
```

Coefficients:

(Intercept)	Solar.R	Wind	Temp
-64.34208	0.05982	-3.33359	1.65209

Holding **all** other variables constant:

- For every 1 unit increase in `Solar.R`, `Ozone` increases by a mean value of 0.06 ppb.
- For every 1 degree increase in `Temp`, `Ozone` increases by a mean value of 1.65 ppb.
- For every 1 unit increase in `Wind`, `Ozone` decreases by a mean value of 3.33 ppb.

## Caution

If the model is not “valid”, then the partial regression coefficients are not meaningful.

# Assumptions

# LINE

As with Simple Linear Regression, we need to check the assumptions of the model (LINE):

- **L**inearity: the relationships between the response and the predictors are all linear.
- **I**ndependence: the observations are independent of each other.
- **N**ormality: the residuals are normally distributed.
- **E**qual variance: the variance of the residuals is constant.

# Recall

In SLR, the model is made up of the **deterministic** component (the line) and the **random** component (the error term).

$$Y_i = \beta_0 + \beta_1 x_i + \epsilon_i$$





**This is the same for MLR:**

$$Y_i = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k + \epsilon_i$$

Since *only* the error term is random, the assumptions are *still* about the error term,  $\hat{\epsilon}$ , which is simple to assess!







# Assumptions of MLR



# Transformation using `log()`

Some evidence of nonlinearity in the diagnostic plots. Transform and re-check assumptions.



# Results

```
Call:
lm(formula = log(Ozone) ~ Solar.R + Wind + Temp, data
= airquality)

Residuals:
    Min       1Q   Median       3Q      Max
-2.06193 -0.29970 -0.00231  0.30756  1.23578

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.2621323   0.5535669   -0.474  0.636798
Solar.R      0.0025152   0.0005567    4.518 1.62e-05 ***
Wind        -0.0615625   0.0157130   -3.918 0.000158 ***
Temp         0.0491711   0.0060875    8.077 1.07e-12 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1

Residual standard error: 0.5086 on 107 degrees of
freedom
(42 observations deleted due to missingness)
```

- All three predictors are statistically significant ( $p < .001$ ).
- The model explains 66% of variance ( $r^2 = 0.66$ ).

# Results compared to SLR

```
Call:
lm(formula = log(Ozone) ~ Solar.R + Wind + Temp, data
= airquality)

Residuals:
    Min       1Q   Median       3Q      Max
-2.06193 -0.29970 -0.00231  0.30756  1.23578

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.2621323   0.5535669  -0.474  0.636798
Solar.R      0.0025152   0.0005567   4.518 1.62e-05 ***
Wind        -0.0615625   0.0157130  -3.918 0.000158 ***
Temp         0.0491711   0.0060875   8.077 1.07e-12 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1

Residual standard error: 0.5086 on 107 degrees of
freedom
(42 observations deleted due to missingness)
```

```
Call:
lm(formula = Ozone ~ Temp, data = airquality)

Residuals:
    Min       1Q   Median       3Q      Max
-40.729 -17.409  -0.587  11.306 118.271

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -146.9955    18.2872  -8.038 9.37e-13 ***
Temp          2.4287     0.2331   10.418 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1

Residual standard error: 23.71 on 114 degrees of
freedom
(37 observations deleted due to missingness)
Multiple R-squared:  0.4877,    Adjusted R-squared:
0.4832
F-statistic: 108.5 on 1 and 114 DF,  p-value: < 2.2e-
```

- All three predictors are statistically significant ( $p < .001$ ).
- The model explains 66% of variance ( $r^2 = 0.66$  vs. 0.48 in SLR).

# Interpretation

# Coefficients

All three predictors are statistically significant ( $p < .001$ ).

- For every 1 unit increase in Solar.R, Log(Ozone) increases by a mean value of 0.0025 ppb, holding all other variables constant.
- For every 1 unit increase in Wind, Log(Ozone) decreases by a mean value of 0.062 ppb, holding all other variables constant.
- For every 1 degree increase in Temp, Log(Ozone) increases by a mean value of 0.049 ppb, holding all other variables constant.

# Residual standard error

On average, the model predicts `Log(Ozone)` within 0.51 ppb of the true value. *Not bad?*

```
[1] 1.665291
```

- On average, the model predicts `Ozone` within 1.6652912 ppb of the true value.
- Number of observations = degrees of freedom (107) + number of parameters in the model (4) = 111.

# R-squared

If there are  $>1$  predictors, use the **Adjusted R-Squared** as it penalises the model for having more predictors that are not useful.



# F-stat

- The F-statistic tests the null hypothesis that all the regression coefficients are equal to zero, i.e.  $H_0 : \beta_1 = \beta_2 = \dots = \beta_k = 0$ .
- As a ratio, it tells us how much better the model is than the null model (i.e. a model with no predictors).
- If the p-value is less than our specified critical value (e.g. 0.05), we reject the null hypothesis and conclude that the current model is better than the null model.

# Reporting

Solar radiation, wind speed and temperature are **significant predictors** of Ozone concentration ( $p < 0.001$ ) with the model accounting for **66% of the variation** in weight.

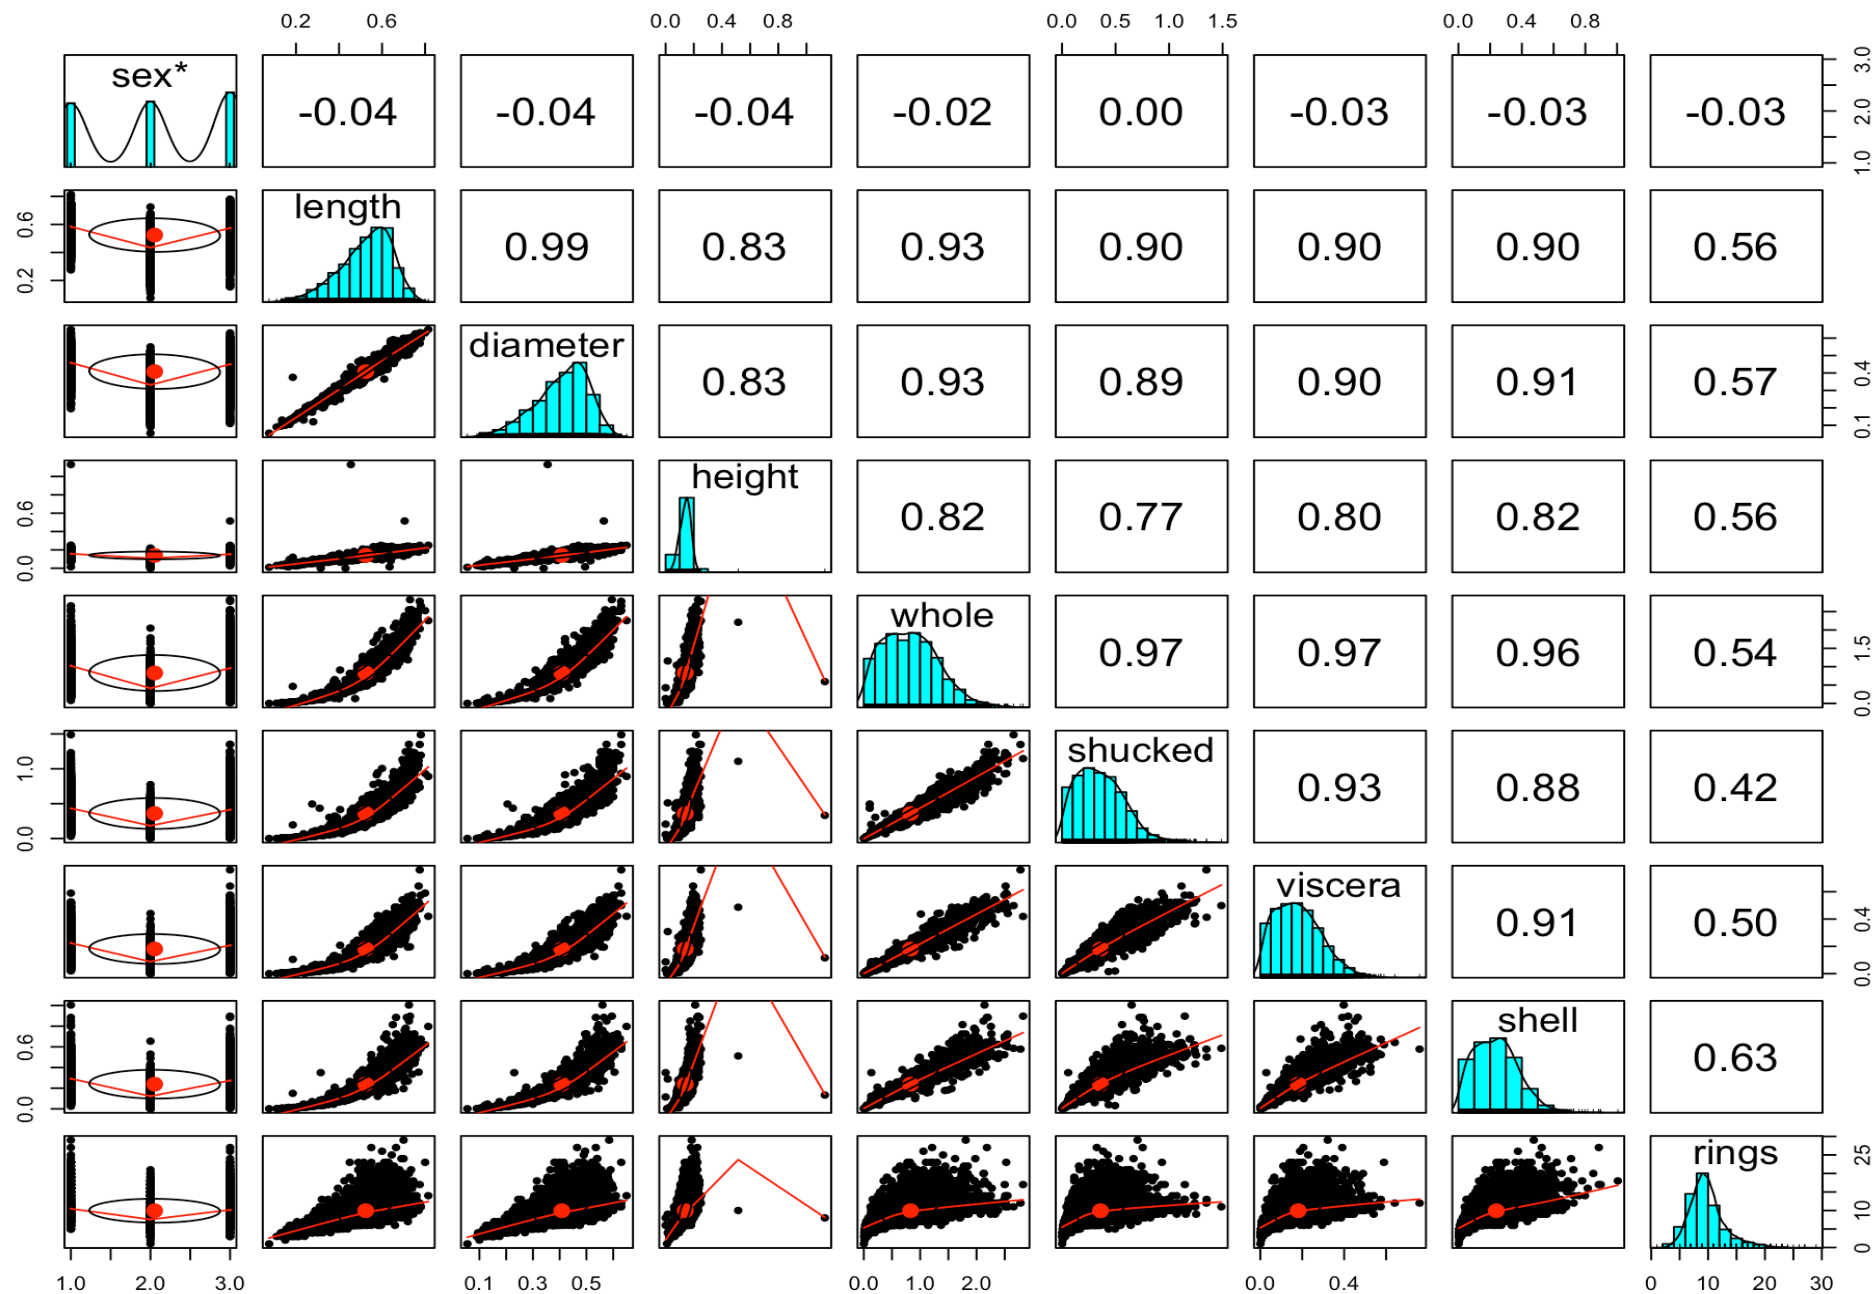
# Abalone: full example

# Data

Data from the [UCI Machine Learning Repository](#).

```
Rows: 4,177
Columns: 9
$ sex      <chr> "M", "M", "F", "M", "I", "I", "F", "F", "M", "F", "F", "M", "...
$ length   <dbl> 0.455, 0.350, 0.530, 0.440, 0.330, 0.425, 0.530, 0.545, 0.475...
$ diameter <dbl> 0.365, 0.265, 0.420, 0.365, 0.255, 0.300, 0.415, 0.425, 0.370...
$ height   <dbl> 0.095, 0.090, 0.135, 0.125, 0.080, 0.095, 0.150, 0.125, 0.125...
$ whole    <dbl> 0.5140, 0.2255, 0.6770, 0.5160, 0.2050, 0.3515, 0.7775, 0.768...
$ shucked  <dbl> 0.2245, 0.0995, 0.2565, 0.2155, 0.0895, 0.1410, 0.2370, 0.294...
$ viscera  <dbl> 0.1010, 0.0485, 0.1415, 0.1140, 0.0395, 0.0775, 0.1415, 0.149...
$ shell    <dbl> 0.150, 0.070, 0.210, 0.155, 0.055, 0.120, 0.330, 0.260, 0.165...
$ rings    <dbl> 15, 7, 9, 10, 7, 8, 20, 16, 9, 19, 14, 10, 11, 10, 10, 12, 7,...
```

# Preview



# Live coding session

Data import → EDA → Model fitting → Diagnostics → Transform/Select → Interpret

*Let's fit a model to predict the whole weight of abalone from other measured variables – I will now switch to RStudio.*

# And we're back!

A quick recap on sub-sampling the dataset:



# What we did

- Fitted a model to predict the whole weight of abalone from other measured variables.
- Performed a transformation of the response variable to improve model fit.
- Checked the assumptions of the model.
- Interpreted the model coefficients.
- Interpreted the model fit.

# Model complexity: overfitting

Why can't we just use ALL the predictors?

# The problem with using too many predictors

- The more predictors you add, the better the model fits the data.
- However, the model may not be able to **generalise** to new data: **overfitting**.

Model	r.squared	adj.r.squared
<code>sqrt(whole) ~ shucked</code>	0.892	0.891
<code>sqrt(whole) ~ shucked + shell</code>	0.952	0.951
<code>sqrt(whole) ~ height + shucked + shell</code>	0.963	0.962
<code>sqrt(whole) ~ length + height + shucked + shell</code>	0.982	0.981
<code>sqrt(whole) ~ length + height + shucked + shell + rings</code>	0.982	0.981
<code>sqrt(whole) ~ length + height + shucked + viscera + shell + rings</code>	0.982	0.981
<code>sqrt(whole) ~ .</code>	0.982	0.981

# The $r^2$ value

The R-squared value is the proportion of variance explained by the model.

$$r^2 = \frac{SS_{reg}}{SS_{tot}} = 1 - \frac{SS_{res}}{SS_{tot}}$$

The adjusted R-squared value is the proportion of variance explained by the model, adjusted for the number of predictors.

$$r_{adj}^2 = 1 - \frac{SS_{res}}{SS_{tot}} \frac{n - 1}{n - p - 1}$$

where  $n$  is the number of observations and  $p$  is the number of predictors.

# Full model vs reduced model

```
Call:
lm(formula = sqrt(whole) ~ ., data = abalone)

Residuals:
    Min       1Q   Median       3Q      Max
-0.218383 -0.016249  0.000771  0.020543  0.105263

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.027849   0.036151  -0.770  0.443065
length       0.959033   0.296239   3.237  0.001678 **
diameter     -0.024686   0.377611  -0.065  0.948019
height       0.969022   0.265067   3.656  0.000427 ***
shucked      0.317776   0.055354   5.741  1.20e-07 ***
viscera      0.107616   0.104461   1.030  0.305614
shell        0.433048   0.095434   4.538  1.72e-05 ***
rings        0.001984   0.001800   1.103  0.273097
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1
```

```
Call:
lm(formula = sqrt(whole) ~ shell + height + diameter,
data = abalone)

Residuals:
    Min       1Q   Median       3Q      Max
-0.149252 -0.030922 -0.004514  0.023821  0.160182

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.13051   0.03854  -3.386  0.001028 **
shell        0.56407   0.09945   5.672  1.49e-07 ***
height       1.33325   0.34613   3.852  0.000212 ***
diameter     1.62282   0.14862  10.919 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1

Residual standard error: 0.05209 on 96 degrees of
freedom
Multiple R-squared:  0.9674,    Adjusted R-squared:
```

- Is the 0.015 improvement in the adjusted R-squared – *an extra 1.5% of the variance explained* – worth the extra predictors?
- Recall: **principle of parsimony** - the simplest model that explains the data is the best.
- But how do we know which predictors to keep?

# Model selection

- Covered in **second year (ENVX2001)**.
- Using techniques of **stepwise regression**, we can select the best model from a set of “candidate” models.
- If we have non-significant predictors, we can consider the effect of removing them from the model (**partial F-test**).
- Aim is to achieve the best balance between **model fit** and **model complexity**.

# Summary

- MLR is an extension of SLR to include more than one predictor.
  - ➡ Instead of a line, we are fitting a “hyperplane” i.e. multiple dimensions.
  - ➡ However, the principles are the same: we are still trying to minimise the sum of squared residuals.
  - ➡ Assumptions of MLR are the same as SLR.
  - ➡ Instead of the multiple R-squared value, we use the adjusted R-squared value to assess model fit.
- Follow the rules of parsimony: the simplest model that explains the data is the best, given similar model fit.
  - ➡ Consider the effect of removing non-significant predictors from the model.

# Thanks!

This presentation is based on the [SOLES Quarto reveal.js template](#) and is licensed under a [Creative Commons Attribution 4.0 International License](#).